

Ming Lin
Staff Algorithm Engineer
DAMO Academy of Alibaba Group
linming04@gmail.com

Alibaba Group (U.S.)
Bellevue
WA 98004
<https://minglin-home.github.io/>

Research Interest

- Research & Development Algorithm Engineer for Cloud-AI and Mobile-AI
- Statistical Machine Learning
- Design state-of-the-art efficient deep learning models
- Real-world applications in computer vision, recommendation system and IoT
- Theoretical foundations for deep learning

Skills & Capabilities

- 14 Years of research and industrial experiences in efficient deep learning, neural architecture design and statistical machine learning
- Programming
 - Python, C/C++
- Deep Learning Frameworks
 - Pytorch, Tensorflow, MXNet

Projects & Achievements

Light-NAS

2021 — 2022

- Homepage [<https://github.com/alibaba/lightweight-neural-architecture-search>]
- Light-NAS is a **full-stack distributed framework for zero-cost NAS**. It is able to design efficient image recognition models and object detection models in 1~2 hours, with state-of-the-art performance on ImageNet, MS COCO, etc.
- Able to optimize multiple inference metrics including **hardware latency** (using latency prediction).
- **Application in Product**
 - Alibaba Pailitao (拍立淘), 10 millions active users per day.
 - Shopping With Your Camera: Visual Image Search Meets E-Commerce at Alibaba (<https://medium.com/coinmonks/shopping-with-your-camera-visual-image-search-meets-e-commerce-at-alibaba-8551925746d0>)

GiraffeDet: High-Performance Object Detection Backbone

2020 — 2021

- Homepage [<https://github.com/damo-cv/GiraffeDet>]

- The GiraffeDet uses an extremely lightweight backbone and a very deep and large neck module which encourages dense information exchange among different spatial scales as well as different levels of latent semantics simultaneously. This design paradigm allows GiraffeDet achieves SOTA performances in object detection tasks with low resource budget.
- **Application in Product**
 - Satellite Remote Sensing Segmentation, 50+ TB high resolution dataset.

Zero-Shot Neural Architecture Search

2020 — 2021

- Homepage [<https://github.com/idstcv/ZenNAS>]
- Zen-NAS is a lightning-fast zero-shot Neural Architecture Search method. Its searching speed is 7800x times faster than EfficientNet while achieving 10x times faster inference speed on NVIDIA T4 GPU with ImageNet top-1 accuracy 83.6%.
- **Used in the winning solution of WebFace260M Track of ICCV21-MFR**, 1.2% improvement over manually designed networks.
- **Application in Product**
 - Dingding M1X FaceID Check-in System, 5+ million devices
 - [Product Homepage] https://tms.dingtalk.com/markets/dingtalk/dingding_M1X_pc

Deep Image Compression

2019 — 2022

- Develop deep learning models for better image compression
- Save 21% bit-rate over BPG

Deep Models for Low Power Devices

2019 — 2020

- Develop and optimize deep neural networks for low power devices (cell-phone, ARM)
- **Used in the winning solution of LPIRC-2019 Competition**

Real-time High-Precision GPS Grid System

2018 — 2019

- Improve GPS precision to 1 centimeter in real-time (1 query per second)
- Positioning successful rate 99.5 % over-all in good weather, 80%-90% for single query.

Early Clinical Diagnosis of Alzheimer's Disease and Major Depressive Disorder

2015 — 2018

- Use large-scale data and machine learning techniques to assist clinical diagnosis of Alzheimer's Disease
- **Able to predict Alzheimer's Disease development in four years**
- Cooperation project with Janssen Research & Development, LCC Michigan Alzheimer's Disease Center

Automated Low-Level Analysis and Description of Diverse Intelligence Video (ALADDIN)

2014 — 2015

- Develop a large-scale content-based video retrieval system
- **Ranked top-1 in 6 out of 8 tasks in MED 2014.**
- Sponsored by Informedia@CMU.

Work Experience

2018 — Present	Staff Algorithm Engineer	DAMO Academy of Alibaba Group (U.S.) <ul style="list-style-type: none"> • Leading a research group of 5 research engineers and several research interns • Efficient deep architecture design • Neural Architecture Search • 8 top conference papers
2015 — 2018	Research Investigator	Department of Computational Medicine and Bioinformatics University of Michigan Medical School Ann Arbor, Michigan 48109, United States. <ul style="list-style-type: none"> • Early diagnosis of Major Depressive Disorder and Alzheimer's disease • Theoretical guarantees for the second order linear models • 2 top-tier conference papers in statistical machine learning • Cooperation with Janssen Research & Development, LCC and Michigan Alzheimer's Disease Center • Mentor: Jieping Ye
2014 — 2015	Postdoctoral Research Fellow	School of Computer Science Carnegie Mellon University Pittsburgh, PA 15213, United States. <ul style="list-style-type: none"> • Adviser: Alexander G. Hauptmann. • Major developer of the ALADDIN project • Video retrieval and feature fusion • 3 top conference papers with 600+ citation in total
2013 — 2014	Visiting Student	School of Computer Science Carnegie Mellon University Pittsburgh, PA 15213, United States. <ul style="list-style-type: none"> • Adviser: Alexander G. Hauptmann.
2012 — 2013	Visiting Student	Department of Computer Science and Engineering Michigan State University

East Lansing, MI 48824, United States.

- Adviser: Rong Jin.

Education

2008 — 2014	Ph.D.	Department of Automation Tsinghua University Beijing 100084, China. <ul style="list-style-type: none">• Adviser: Changshui Zhang
2004 — 2008	Bachelor	Department of Automation Tsinghua University Beijing 100084, China.

Peer-Review Service

- Reviewer of AAAI, IJCAI, ICML, NIPS, ICLR, CVPR, ICCV etc.
- Reviewer of Journals:
 - o IEEE Transactions on Pattern Analysis and Machine Intelligence
 - o ACM Transactions on Knowledge Discovery from Data
 - o Pattern Recognition
 - o Neurocomputing
 - o Computer Vision and Image Understanding.
 - o Data Mining and Knowledge Discovery

Impact

Google citation 1034, h-index 13, i10-index 16

Selected Publications & Pre-prints

- Ming Lin, Pichao Wang, Zhenhong Sun, Heseng Chen, Xiuyu Sun, Qi Qian, Hao Li, Rong Jin. **Zen-NAS: A Zero-Shot NAS for High-Performance Deep Image Recognition**. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), 2021.
- Zhenhong Sun, Ming Lin, Xiuyu Sun, Zhiyu Tan, Rong Jin. **ZenDet: Revisiting Efficient Object Detection Backbones from Zero-Shot Neural Architecture Search**. arXiv 2021.
- Zhenzhong Lan, Ming Lin, Xuanchong Li, Alexander G. Hauptmann, Bhiksha Raj. **Beyond Gaussian Pyramid: Multi-skip Feature Stacking for Action Recognition**. IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Pages 204-212, 2015.

Bibliography

Peer-Reviewed Conferences

1. Yichen Qian, Xiuyu Sun, Ming Lin, Zhiyu Tan, Rong Jin. Entroformer: A Transformer-based Entropy Model for Learned Image Compression. In Proceedings of the International Conference on Learning Representations (ICLR), 2022.
2. Yaohua Wang, Yaobin Zhang, Fangyi Zhang, Senzhang Wang, Ming Lin, YuQi Zhang, Xiuyu Sun. Ada-NETS: Face Clustering via Adaptive Neighbour Discovery in the Structure Space. In Proceedings of the International Conference on Learning Representations (ICLR), 2022.
3. Yiqi Jiang, Zhiyu Tan, Junyan Wang, Xiuyu Sun, Ming Lin, Hao Li. GiraffeDet: A Heavy-Neck Paradigm for Object Detection. In Proceedings of the International Conference on Learning Representations (ICLR), 2022.
4. Ming Lin, Pichao Wang, Zhenhong Sun, Heseng Chen, Xiuyu Sun, Qi Qian, Hao Li, Rong Jin. Zen-NAS: A Zero-Shot NAS for High-Performance Deep Image Recognition. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), 2021.
5. Yichen Qian, Zhiyu Tan, Xiuyu Sun, Ming Lin, Dongyang Li, Zhenhong Sun, Li Hao, Rong Jin. Learning Accurate Entropy Model with Global Reference for Image Compression. In Proceedings of the International Conference on Learning Representations (ICLR), 2021.
6. Heseng Chen, Ming Lin, Xiuyu Sun, Qian Qi, Hao Li, Rong Jin. MuffNet: Multi-Layer Feature Federation for Mobile Deep Learning. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV workshop), 2019.
7. Ming Lin, Xiaomin Song, Qi Qian, Hao Li, Liang Sun, Shenghuo Zhu, Rong Jin. Robust Gaussian Process Regression for Real-Time High Precision GPS Signal Enhancement. In Proceedings of the 25TH ACM SIGKDD CONFERENCE ON KNOWLEDGE DISCOVERY AND DATA MINING (SIGKDD), 2019.
8. Ming Lin, Shuang Qiu, Jieping Ye, Xiaomin Song, Qi Qian, Liang Sun, Shenghuo Zhu, Rong Jin. Which Factorization Machine Modeling is Better: A Theoretical Answer with Optimal Guarantee. The Thirty-Third AAAI Conference on Artificial Intelligence (AAAI), 2019.
9. Tieliang Gong, Guangtao Wang, Jieping Ye, Zongben Xu, Ming Lin. Margin Based PU Learning. In Proceedings of the 32nd AAAI Conference on Artificial Intelligence (AAAI), 2018.
10. Xiang Li, Aoxiao Zhong, Ming Lin, Ning Guo, Mu Sun, Arkadiusz Sitek, Jieping Ye, James Thrall, Quanzheng Li. Self-paced Convolutional Neural Network for Computer Aided Detection in Medical Imaging Analysis. In Proceedings of the 8th International Workshop on Machine Learning in Medical Imaging (MLMI), 2017.
11. Zhenzhong Lan, Shou-I Yu, Dezhong Yao, Ming Lin, Bhiksha Raj ; Alexander Hauptmann. The Best of Both Worlds: Combining Data-Independent and Data-Driven Approaches for Action Recognition. IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW). Pages 1196-1205, 2016.
12. Ming Lin, Jieping Ye. A Non-convex One-Pass Framework for Generalized Factorization Machine and Rank-One Matrix Sensing. In Proceedings of the 30th Annual Conference on Neural Information Processing Systems (NIPS), Pages 1633-1641, 2016.
13. Chuang Gang, Ming Lin, Yi Yang, Gerard de Melo, Alexander G. Hauptmann. Concepts Not Alone: Exploring Pairwise Relationships for Zero-Shot Video Activity Recognition. In

Proceedings of the 30th AAAI Conference on Artificial Intelligence (AAAI), Pages 3487-3493, 2016.

14. Chuang Gang, Ming Lin, Yi Yang, Alexander G. Hauptmann. Exploring Semantic Inter-Class Relationships (SIR) for Zero-Shot Action Recognition. In Proceedings of the 29th AAAI Conference on Artificial Intelligence (AAAI), Pages 3769-3775, 2015
15. Zhenzhong Lan, Ming Lin, Xuanchong Li, Alexander G. Hauptmann, Bhiksha Raj. Beyond Gaussian Pyramid: Multi-skip Feature Stacking for Action Recognition. IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Pages 204-212, 2015.
16. Ming Lin, Zhenzhong Lan, Alexander G. Hauptmann. Density Corrected Sparse Recovery when R.I.P. Condition is Broken. In Proceedings of the 24th International Joint Conference on Artificial Intelligence (IJCAI), Pages 3664-3670, 2015.
17. Shou-I Yu, Lu Jiang, Zexi Mao, Xiaojun Chang, Xingzhong Du, Chuang Gan, Zhenzhong Lan, Zhongwen Xu, Xuanchong Li, Yang Cai, Anurag Kumar, Yajie Miao, Lara Martin, Nikolas Wolfe, Shicheng Xu, Huan Li, Ming Lin, Zhigang Ma, Yi Yang, Deyu Meng, Shiguang Shan, Pinar Duygulu Sahin, Susanne Burger, Florian Metze, Rita Singh, Bhiksha Raj, Teruko Mitamura, Richard Stern, Alexander Hauptmann. Informedia@ TRECVID 2014 MED and MER. NIST TRECVID Video Retrieval Evaluation Workshop, 2014.
18. Ming Lin, Rong Jin, Changshui Zhang. Efficient Sparse Recovery via Adaptive Non-Convex Regularizers with Oracle Property. In Proceedings of the 30th Conference on Uncertainty in Artificial Intelligence (UAI), Pages 505-514, 2014.
19. Lijun Zhang, Jinfeng Yi, Ming Lin, Xiaofei He. Online Kernel Learning with a Near Optimal Sparsity Bound. In Proceedings of the 30th International Conference on Machine Learning (ICML), Pages 621 – 629, 2013.

Peer-Reviewed Journals

1. Jian Liang, Kun Chen, Ming Lin, Changshui Zhang, Fei Wang. Robust Finite Mixture Regression for Heterogeneous Targets. Data Mining and Knowledge Discovery, Volume 32, Issue 6, pp 1509–1560, November 2018.
2. Ming Lin, Pinghua Gong, Tao Yang, Jieping Ye, Roger L. Albin, Hiroko H. Dodge. Big Data Analytical Approaches to the NACC Dataset: Aiding Preclinical Trial Enrichment, Volume 32, Issue1, Pages 18-27, January-March 2018.
3. Daqing Chang, Ming Lin, Changshui Zhang. On the generalization ability of online gradient descent algorithm under the quadratic growth condition. To appear in TNNLS 2018.
4. Ming Lin, Vaibhav Narayan, Wayne C. Drevets, Jieping Ye, Qingqin Li. Application of Growth Mixture Modeling in Antidepressant Treatment Response Studies. Biological Psychiatry, Volume 81, Issue 10, Supplement, Page S224, May 2017.
5. Xiaojun Chang, Zhigang Ma, Ming Lin, Yi Yang, Alexander G. Hauptmann. Feature Interaction Augmented Sparse Learning for Fast Kinect Motion Detection. IEEE Transactions on Image Processing, Volume 26, Issue 8, Pages 3911-3920. 2017.
6. Shou-I Yu, Yi Yang, Zhongwen Xu, Shicheng Xu, Deyu Meng, Zexi Mao, Zhigang Ma, Ming Lin, Xuanchong Li, Huan Li, Zhenzhong Lan, Lu Jiang, Alexander G. Hauptmann, Chuang Gan, Xingzhong Du, Xiaojun Chang. Strategies for Searching Video Content with Text Queries or Video Examples (Invited Paper). IEEE Transactions on Media Technology and Applications 4.3, Pages 227-238, 2016

7. Ming Lin, Lijun Zhang, Rong Jin, Shifeng Weng, Changshui Zhang. Online Kernel Learning with Nearly Constant Support Vectors. *Neurocomputing*. Volume 179, Pages 26-36, 2016.
8. Zhen Hu, Ming Lin, Changshui Zhang. Dependent Online Kernel Learning with Constant Number of Random Fourier Features. *IEEE Transactions on Neural Networks and Learning Systems (TNNLS)*, Volume 26, Issue 10, Pages 2464-2476, 2015.
9. Zheng Pan, Ming Lin, Guangdong Hou, Changshui Zhang. Damping proximal coordinate descent algorithm for non-convex regularization. *Neurocomputing*, vol 152, Pages 151-163, 2015.
10. Ming Lin, Fei Wang, Changshui Zhang. Large-Scale Eigenvector Approximation via Hilbert Space Embedding Nystrom. *Pattern Recognition (PR)*, 48(5), Pages 1904-1912, 2015.
11. Ming Lin, Shifeng Weng, Changshui Zhang. On the Sample Complexity of Random Fourier Features for Online Learning. *ACM Transactions on Knowledge Discovery from Data (TKDD)*, Volume 8 Issue 3, Pages 13:1--13:19, June 2014.
12. Shizhun Yang, Ming Lin, Chenping Hou, Changshui Zhang, Yi Wu. A General Framework for Transfer Sparse Subspace Learning. *Neural Computing and Applications*. Volume 21, Number 7, Pages 1801-1817, August 2012.